

# Assignment-2 Solutions

Question 1 :

In PyTorch , we calculate gradients using a dynamic computational graph.

- Behaviour of loss.backward() : when we call this method,Pytorch computes the gradient of the loss with respect to all leaf tensors that have `requires_grad=True`.These values are stored in the `.grad` which is an attribute of each tensor.
- Pytorch adds the newly calculated gradients to the existing values in the `.grad` attribute rather than over writing them.
- The requirement of optimizer.zero\_grad() : Because the gradients accumulate , if we do not clear them at the start of every training step,the gradients from the previous batch will be added to the current ones.

If forgotten the update rule becomes

$$\theta_{new} = \theta_{old} - \eta \cdot \sum_{i=1}^{current\_step} \nabla Loss_i$$

The above formula causes the gradients to grow incorrectly , leading to divergent training or incorrect weight updates.

Question 2:

`.view()`:

It requires the tensor to be contiguous in memory. It always returns a "view" of the original data(no copying).It will give an error if the tensor is non contiguous like after a transpose

`.reshape()`

It can handle both contiguous and non contiguous tensors. It returns a view if possible , otherwise it copies the data to a new memory block.We will have safer fall back because it handles the memory layout automatically.

Question-3:

If we attempt an operation(like addition) between a tensor on the CPU and a tensor on CUDA:0,Pytorch will then raise a runtime error.All tensors involved in a single operation must reside on the same device.

Tensor Creation in forward : Moving a model to the GPU via `model.to('cuda')` only moves the model's existing parameters and buffers. Any new tensors created inside the forward method like example `torch.zeros(5)` will be created on the CPU by default.You must manually assign them to the correct device using `device=x.device` or `.to(device)`.