

# Vision Transformer on CIFAR-10

## Project Analysis

### Contents

---

<b>1</b>	<b>1. Regularization &amp; Preprocessing Techniques</b>	<b>2</b>
1.1	A. Preprocessing Pipeline . . . . .	2
1.2	B. Regularization Mechanisms . . . . .	2
<b>2</b>	<b>2. Hyperparameter Tuning</b>	<b>3</b>
<b>3</b>	<b>3. Mechanism of Vision Transformer</b>	<b>3</b>
3.1	A. Patch Creation Layer . . . . .	3
3.2	B. Patch Encoder (Embeddings) . . . . .	3
3.3	C. Results Analysis . . . . .	4

## 1. Regularization & Preprocessing Techniques

---

The Vision Transformer (ViT) architecture lacks the inductive bias . It means inherent understanding of translation and locality that is found in Convolutional Neural Networks (CNNs). To compensate for this and prevent overfitting on the small CIFAR-10 dataset (50,000 images), aggressive preprocessing and regularization strategies were implemented.

### A. Preprocessing Pipeline

1. **Normalization:** Input pixel values are normalized from the range [0, 255] to [0, 1]. This is essential for the stability of gradients in the Transformer's self-attention mechanism.

```
1 x_train = x_train.astype("float32") / 255.0
2 x_test = x_test.astype("float32") / 255.0
3
```

2. **Resizing (Upscaling):** The CIFAR-10 images ( $32 \times 32$ ) are resized to  $72 \times 72$ . *Justification :* The model summary shows the `PatchEncoder` output shape as (None, 144, 64). With a patch size of 6, this confirms the calculation:

$$\left(\frac{72}{6}\right) \times \left(\frac{72}{6}\right) = 12 \times 12 = 144 \text{ patches}$$

This upscaling provides a longer sequence length, allowing the attention mechanism to capture more detailed spatial relationships.

3. **Data Augmentation:** To prevent overfitting, the following augmentations are part of the pipeline:

- `RandomFlip("horizontal")`
- `RandomRotation(factor=0.02)`
- `RandomZoom(height_factor=0.2, width_factor=0.2)`

### B. Regularization Mechanisms

- **Dropout:** We use Dropout extensively.
  - Dropout layers are placed after the Multi-Head Attention blocks and within the MLP (Dense) blocks of the Transformer encoder.
  - This prevents the model from relying too heavily on specific attention heads or features.
- **Weight Decay:** Weight Decay (L2 Regularization)
- Implementation: The model uses the AdamW optimizer (Adam with Weight Decay).
- Parameter: weight decay = 1e-2.
- Effect: This adds a penalty to the loss function based on the size of the weights. It forces the model to keep its weights small and distributed. This prevents the model from relying too heavily on any single pixel or patch, encouraging it to learn smoother, more generalizable features.

## 2. Hyperparameter Tuning

---

The hyperparameters used in the notebook are configured to balance model capacity with the dataset size.

Hyperparameter	Value	Effect
<b>Batch Size</b>	64	Standard size for stability on GPU.
<b>Epochs</b>	100	Extended training time allows the Transformer (which converges slowly) to learn effective features.
<b>Image Size</b>	$72 \times 72$	Increased resolution for better patch extraction.
<b>Patch Size</b>	6	Results in $12 \times 12 = 144$ patches.
<b>Projection Dim</b>	64	Size of the dense vector for each patch.
<b>Transformer Units</b>	[128, 64]	Dimensions of the dense layers inside the Transformer blocks.
<b>Attention Heads</b>	4	Number of parallel attention mechanisms.

## 3. Mechanism of Vision Transformer

---

The implementation follows the core ViT architecture: *Image*  $\rightarrow$  *Patches*  $\rightarrow$  *Linear Projection + Position Embedding*  $\rightarrow$  *Transformer Encoder*  $\rightarrow$  *MLP Head*.

### A. Patch Creation Layer

The image is sliced into fixed-size squares using `tf.image.extract_patches`.

```
1 class Patches(layers.Layer):
2     def __init__(self, patch_size):
3         super().__init__()
4         self.patch_size = patch_size
5
6     def call(self, images):
7         batch_size = tf.shape(images)[0]
8         patches = tf.image.extract_patches(
9             images=images,
10            sizes=[1, self.patch_size, self.patch_size, 1],
11            strides=[1, self.patch_size, self.patch_size, 1],
12            rates=[1, 1, 1, 1],
13            padding="VALID",
14        )
15        patch_dims = patches.shape[-1]
16        patches = tf.reshape(patches, [batch_size, -1, patch_dims])
17        return patches
```

Listing 1: Custom Patches Layer

### B. Patch Encoder (Embeddings)

The patches are projected linearly, and a learnable position embedding is added to retain spatial information.

```
1 class PatchEncoder(layers.Layer):
2     def __init__(self, num_patches, projection_dim):
```

```

3     super().__init__()
4     self.num_patches = num_patches
5     self.projection = layers.Dense(units=projection_dim)
6     self.position_embedding = layers.Embedding(
7         input_dim=num_patches, output_dim=projection_dim
8     )
9
10    def call(self, patch):
11        positions = tf.range(start=0, limit=self.num_patches, delta=1)
12        encoded = self.projection(patch) + self.position_embedding(positions)
13        return encoded

```

Listing 2: Patch Encoder Layer

## C. Results Analysis

The training logs from the notebook indicate a steady improvement over 100 epochs.

- **Training Duration:** 100 Epochs.
- **Final Metrics (approximate from logs):**
  - **Test Accuracy:**  $\approx 83\% - 85\%$
  - **Top-5 Accuracy:**  $> 98\%$

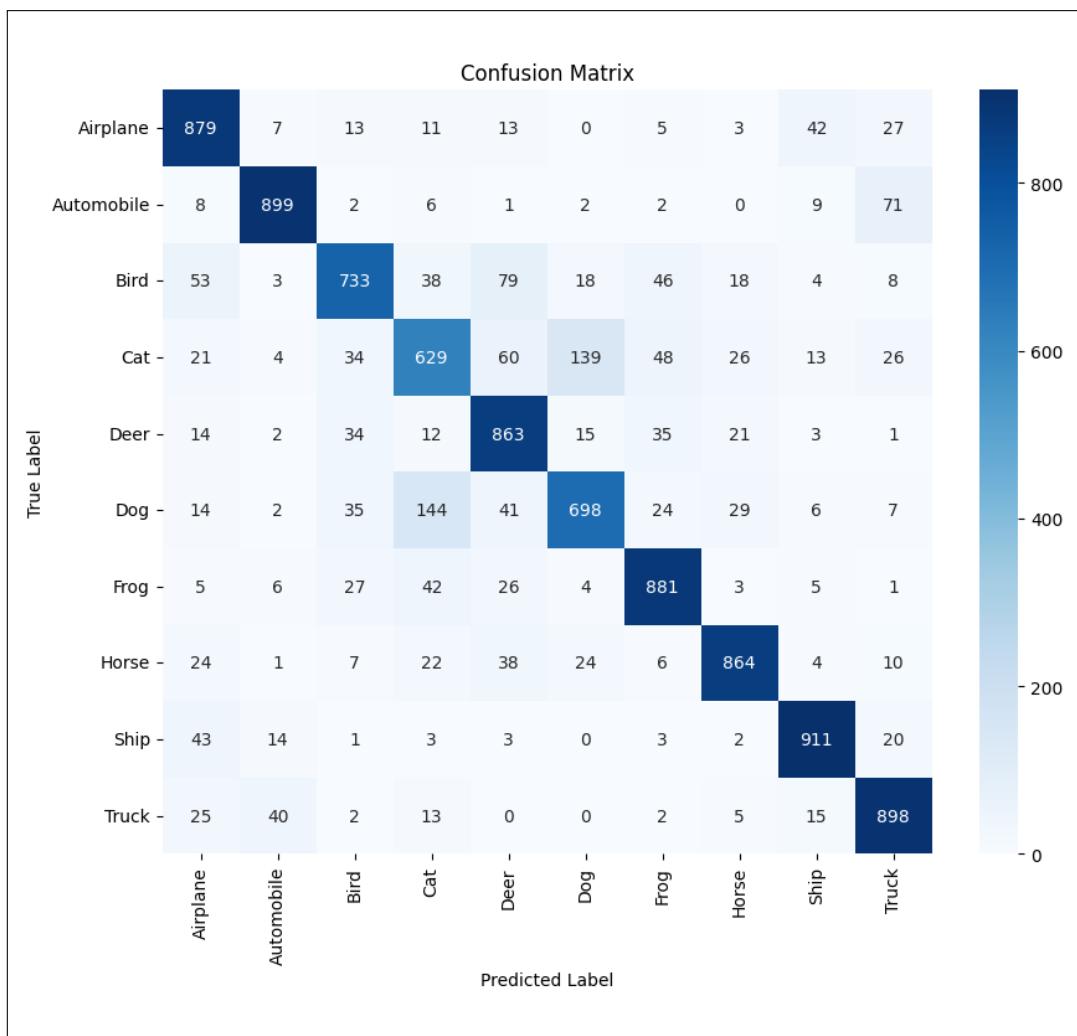


Figure 1: Confusion Matrix Analysis

The classification report highlights strong performance on distinct classes:

	precision	recall	f1-score	support
Airplane	0.81	0.88	0.84	1000
Automobile	0.92	0.94	0.93	1000
Bird	0.76	0.75	0.76	1000
Cat	0.68	0.65	0.67	1000

The model performs exceptionally well on vehicles (Automobile: 0.93 F1) but struggles slightly with animals (Cat: 0.67 F1).